



Final Round Presentation

Himank Gupta

Problem Statement: Your job is to design a next-generation Claims Processing product for vehicle insurance companies that leverages agentic AI benefits.

Status Quo: What is the problem? Why is it a problem?

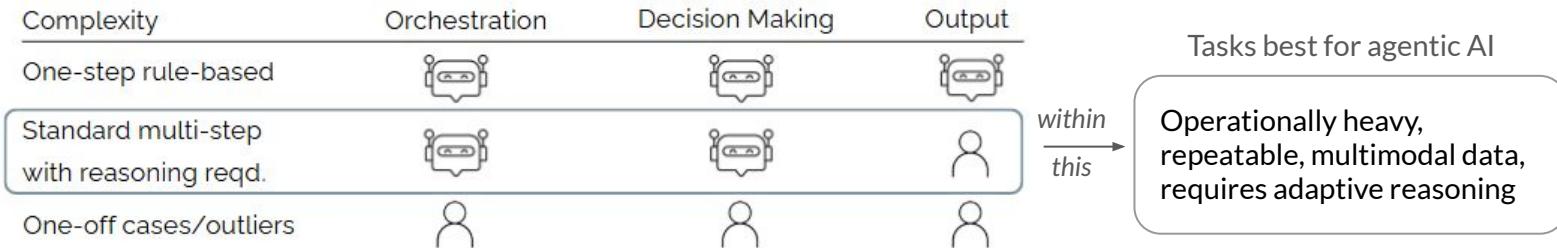
1. **Vehicle insurance claims is a fragmented, error-prone, manual workflow**
 - a. FNOL -> Triage and Assignment -> Investigation and Adjudication -> Settlement and Payout
2. **All stakeholders struggle, leaving much to be desired**

Stakeholder	Pain Points	Desires
Claimant	Confused, anxious, no visibility on progress	Reassurance & quick, transparent updates
Employee	Tedious, repetitive work, fragmented data	Less chaos, more stimulating tasks
Insurance co.	Sink of resources, poor customer experience	Accurate, efficient claims handling process

3. **There is a significant impact^{1,2,3} of continuing with outdated processes**
 - a. **40%** of adjuster time (and **30%** of premium spent) on administrative tasks instead of decision-making
 - b. **38%** of customers fall into the lowest satisfaction tier; **30%** higher retention potential w/ 24hr closure
 - c. **80%** cost reduction potential by moving to automated claims handling: **\$50Bn+** opportunity over 5yrs

Opportunity: What should we solve? What's our goal?

1. Certain types of claims and tasks are more suited for disruption (agentic AI automation)



2. By improving the process, we can move two critical metrics that will directly drive business outcomes

% reduction in cost per claim

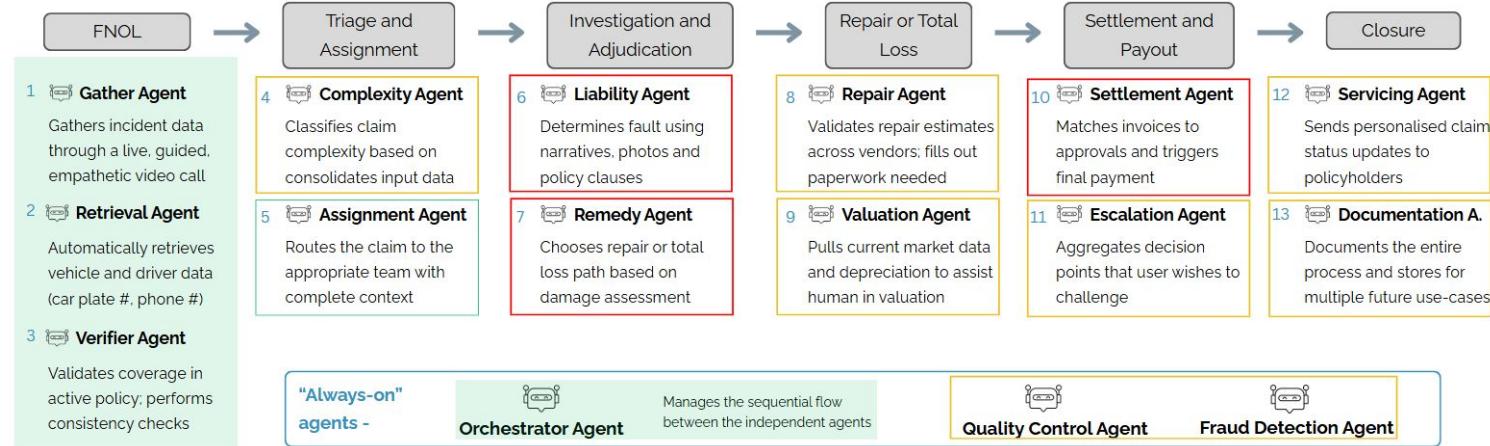
(time, inaccuracies)

Customer satisfaction score (CSAT)

(quicker, easier)

Options: What can be built? What to prioritise?

1. There are several agentic opportunities across each stage of the claims process

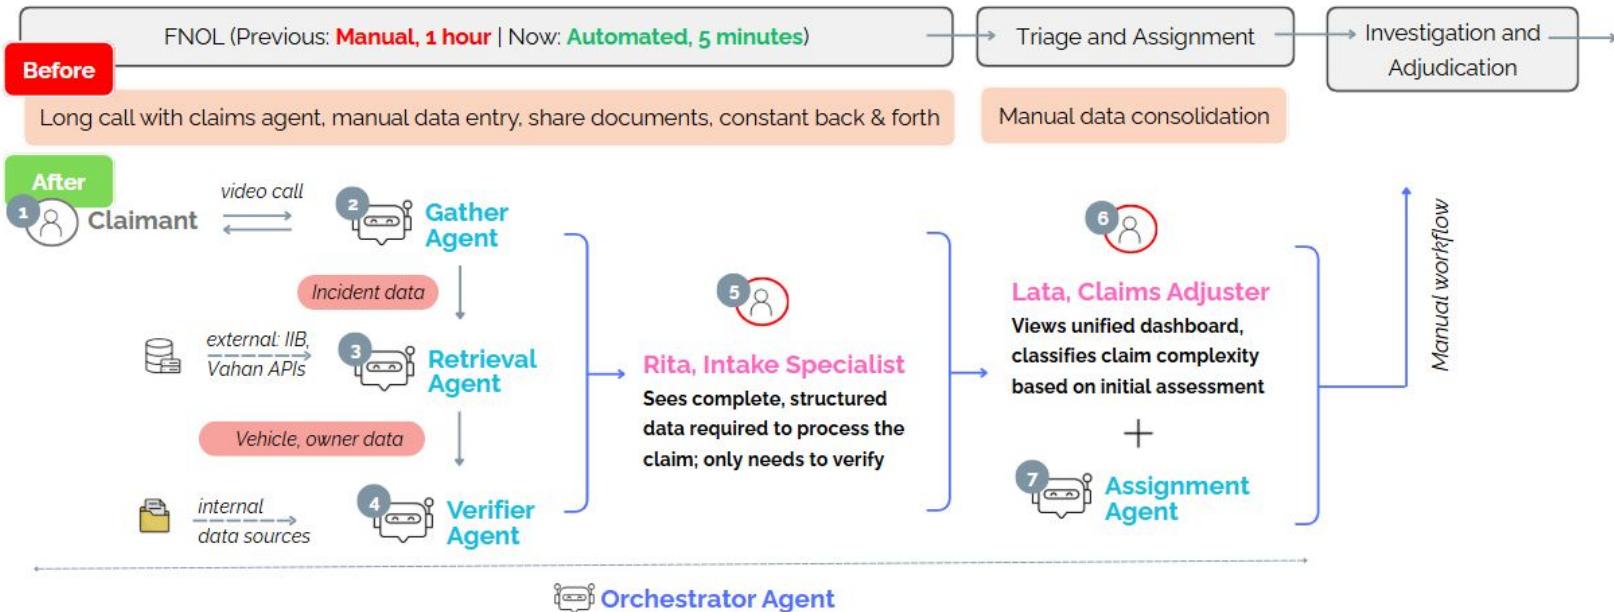


2. The MVP objective is reduced \$/claim & higher CSAT with practicality in mind | Key prioritisation metrics:

- Efficiency gain (\downarrow hours, \uparrow quality) - Redundant, repetitive data aggregation and validation
- Auditability of AI engine (ease in catching mistakes) - Clearly traceable steps with low subjectivity
- Effort (\downarrow hours for build and integration) - Modular, minimal conflicts with legacy workflow
- Risk if bad output (\$ loss, resources to fix) - Minimal regulatory implications and financial involvement
- Customer cognitive load (time spent interacting) - Reduced touchpoints, increased internal data pulling

Introducing 'Arete': What's the improved workflow in MVP?

1. The new workflow will have agentic AI-driven data collection, validation and assignment:



2. Artefacts: MVP ([Working prototype](#)) to illustrate the user journey, Retrieval Agent (experimental, out of scope) programmed a v1 of one functionality for the agent

Metrics and Guardrails: What to track? What to keep in mind?

1. Key metrics to be tracked to help us understand whether our product is achieving the objective:

Metric Type	Metrics
Business	% cost reduction per claim, CSAT, customer retention rate
Product NSM	average time to process a claim, % monthly claims successfully processed without errors
Agentic	time spent by user on submitting a claim, % claim_data objects approved by intake specialist
Counter	# data issue tickets generated downstream, average human re-work time spent per claim

2. Guardrails in MVP:

- Human-in-the-loop - AI is instructed to redirect to human if confused; human override option is 24/7
- Identity verification - Verify the legitimacy of claimant with increasing autonomous workflows

3. Risks and mitigation:

- Technology risk - This plan rests on engineering excellence and innovation capability; planning is crucial
- Customer acceptance - Serious, concerning event, don't want to be 'tests'; quality + edge cases > speed